

```
# Data analysis and visualization
import tensorflow as tf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

# Preprocessing and evaluation
from sklearn.model_selection import train_test_split
from sklearn.compose import make_column_transformer
from sklearn.preprocessing import MinMaxScaler
```

```
(X_train , y_train), (X_test , y_test) = tf.keras.datasets.boston_housing.load_data(
    path = 'boston_housing.npz',
    test_split = 0.2,
    seed = 42
)
```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston_housing.npz
57026/57026 [=====] - 0s 0us/step

```
# Checking the data shape and type
(X_train.shape, type(X_train)), (X_test.shape, type(X_test)), (y_train.shape, type(y_train)), (y_test.shape, type(y_test))

(((404, 13), numpy.ndarray),
 ((102, 13), numpy.ndarray),
 ((404,), numpy.ndarray),
 ((102,), numpy.ndarray))
```

```
# Converting Data to DataFrame
X_train_df = pd.DataFrame(X_train)
y_train_df = pd.DataFrame(y_train)
```

```
# Preview the training data
X_train_df.head(10)
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.09178	0.0	4.05	0.0	0.510	6.416	84.1	2.6463	5.0	296.0	16.6	395.50	9.04
1	0.05644	40.0	6.41	1.0	0.447	6.758	32.9	4.0776	4.0	254.0	17.6	396.90	3.53
2	0.10574	0.0	27.74	0.0	0.609	5.983	98.8	1.8681	4.0	711.0	20.1	390.11	18.07
3	0.09164	0.0	10.81	0.0	0.413	6.065	7.8	5.2873	4.0	305.0	19.2	390.91	5.52
4	5.09017	0.0	18.10	0.0	0.713	6.297	91.8	2.3682	24.0	666.0	20.2	385.09	17.27
5	0.10153	0.0	12.83	0.0	0.437	6.279	74.5	4.0522	5.0	398.0	18.7	373.66	11.97
6	0.31827	0.0	9.90	0.0	0.544	5.914	83.2	3.9986	4.0	304.0	18.4	390.70	18.33
7	0.29090	0.0	21.89	0.0	0.624	6.174	93.6	1.6119	4.0	437.0	21.2	388.08	24.16
8	4.03841	0.0	18.10	0.0	0.532	6.229	90.7	3.0993	24.0	666.0	20.2	395.33	12.87
9	0.22438	0.0	9.69	0.0	0.585	6.027	79.7	2.4982	6.0	391.0	19.2	396.90	14.33

```
# View summary of datasets
X_train_df.info()
print('_'*40)
y_train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 404 entries, 0 to 403
Data columns (total 13 columns):
#   Column  Non-Null Count  Dtype
---  -
0     0      404 non-null     float64
1     1      404 non-null     float64
2     2      404 non-null     float64
3     3      404 non-null     float64
4     4      404 non-null     float64
5     5      404 non-null     float64
6     6      404 non-null     float64
7     7      404 non-null     float64
8     8      404 non-null     float64
9     9      404 non-null     float64
10    10      404 non-null     float64
11    11      404 non-null     float64
```

```
12 12      404 non-null    float64
dtypes: float64(13)
memory usage: 41.2 KB

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 404 entries, 0 to 403
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      404 non-null    float64
dtypes: float64(1)
memory usage: 3.3 KB
```

```
# distribution of numerical feature values across the samples
X_train_df.describe()
```

	0	1	2	3	4	5	6	7	8	9	10
count	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000
mean	3.789989	11.568069	11.214059	0.069307	0.554524	6.284824	69.119307	3.792258	9.660891	408.960396	18.481931
std	9.132761	24.269648	6.925462	0.254290	0.116408	0.723759	28.034606	2.142651	8.736073	169.685166	2.157322
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.137000	1.000000	187.000000	12.600000
25%	0.081960	0.000000	5.190000	0.000000	0.452000	5.878750	45.475000	2.097050	4.000000	281.000000	17.400000
50%	0.262660	0.000000	9.690000	0.000000	0.538000	6.210000	77.500000	3.167500	5.000000	330.000000	19.100000
75%	3.717875	12.500000	18.100000	0.000000	0.624000	6.620500	94.425000	5.118000	24.000000	666.000000	20.200000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000

```
# Create column transformer
ct = make_column_transformer(
    (MinMaxScaler(), [0, 1, 2, 4, 5, 6, 7, 8, 9, 10, 11, 12])
)

# Normalization and data type change
X_train = ct.fit_transform(X_train).astype('float32')
X_test = ct.transform(X_test).astype('float32')
y_train = y_train.astype('float32')
y_test = y_test.astype('float32')

# Distribution of X_train feature values after normalization
pd.DataFrame(X_train).describe()
```

	0	1	2	3	4	5	6	7	8	9	10
count	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000
mean	0.042528	0.115681	0.394210	0.348815	0.521905	0.681970	0.241618	0.376560	0.423589	0.625737	0.897607
std	0.102650	0.242696	0.253866	0.239522	0.138678	0.288719	0.194973	0.379829	0.323827	0.229502	0.232131
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000850	0.000000	0.173387	0.137860	0.444098	0.438466	0.087361	0.130435	0.179389	0.510638	0.944992
50%	0.002881	0.000000	0.338343	0.314815	0.507569	0.768280	0.184767	0.173913	0.272901	0.691489	0.985892
75%	0.041717	0.125000	0.646628	0.491770	0.586223	0.942585	0.362255	1.000000	0.914122	0.808511	0.997252
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

```
# Reserve data for validation
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.1, random_state=42)
X_train.shape, X_val.shape, y_train.shape, y_val.shape

((363, 12), (41, 12), (363,), (41,))
```

```

# Set random seed
tf.random.set_seed(42)

# Building the model
model = tf.keras.Sequential([
    tf.keras.layers.Dense(units=10, activation='relu', input_shape=(X_train.shape[1],), name='Dense_1'),
    tf.keras.layers.Dense(units=100, activation='relu', name='Dense_2'),
    tf.keras.layers.Dense(units=1, name='Prediction')
])

# Compiling the model
model.compile(
    loss = tf.keras.losses.mean_squared_error,
    optimizer = tf.keras.optimizers.RMSprop(learning_rate=0.01),
    metrics = ['mse']
)

# Training the model
history = model.fit(
    X_train,
    y_train,
    batch_size=32,
    epochs=50,
    validation_data=(X_val, y_val)
)

Epoch 1/50
12/12 [=====] - 1s 33ms/step - loss: 308.4484 - mse: 308.4484 - val_loss: 146.9222 - val_mse: 146.9222
Epoch 2/50
12/12 [=====] - 0s 9ms/step - loss: 95.1682 - mse: 95.1682 - val_loss: 92.9939 - val_mse: 92.9939
Epoch 3/50
12/12 [=====] - 0s 7ms/step - loss: 62.2104 - mse: 62.2104 - val_loss: 68.8015 - val_mse: 68.8015
Epoch 4/50
12/12 [=====] - 0s 8ms/step - loss: 49.0687 - mse: 49.0687 - val_loss: 56.8894 - val_mse: 56.8894
Epoch 5/50
12/12 [=====] - 0s 8ms/step - loss: 45.6809 - mse: 45.6809 - val_loss: 68.7085 - val_mse: 68.7085
Epoch 6/50
12/12 [=====] - 0s 7ms/step - loss: 38.9542 - mse: 38.9542 - val_loss: 61.3664 - val_mse: 61.3664
Epoch 7/50
12/12 [=====] - 0s 9ms/step - loss: 32.9990 - mse: 32.9990 - val_loss: 35.6852 - val_mse: 35.6852
Epoch 8/50
12/12 [=====] - 0s 10ms/step - loss: 29.6053 - mse: 29.6053 - val_loss: 47.0410 - val_mse: 47.0410
Epoch 9/50
12/12 [=====] - 0s 8ms/step - loss: 28.8039 - mse: 28.8039 - val_loss: 30.4692 - val_mse: 30.4692
Epoch 10/50
12/12 [=====] - 0s 8ms/step - loss: 28.7124 - mse: 28.7124 - val_loss: 38.0187 - val_mse: 38.0187
Epoch 11/50
12/12 [=====] - 0s 5ms/step - loss: 23.2083 - mse: 23.2083 - val_loss: 45.0314 - val_mse: 45.0314
Epoch 12/50
12/12 [=====] - 0s 5ms/step - loss: 26.8295 - mse: 26.8295 - val_loss: 25.9321 - val_mse: 25.9321
Epoch 13/50
12/12 [=====] - 0s 5ms/step - loss: 22.4063 - mse: 22.4063 - val_loss: 33.7716 - val_mse: 33.7716
Epoch 14/50
12/12 [=====] - 0s 5ms/step - loss: 23.0061 - mse: 23.0061 - val_loss: 19.8987 - val_mse: 19.8987
Epoch 15/50
12/12 [=====] - 0s 6ms/step - loss: 22.7476 - mse: 22.7476 - val_loss: 19.6241 - val_mse: 19.6241
Epoch 16/50
12/12 [=====] - 0s 5ms/step - loss: 21.5558 - mse: 21.5558 - val_loss: 22.3366 - val_mse: 22.3366
Epoch 17/50
12/12 [=====] - 0s 6ms/step - loss: 22.3931 - mse: 22.3931 - val_loss: 46.4523 - val_mse: 46.4523
Epoch 18/50
12/12 [=====] - 0s 7ms/step - loss: 23.5593 - mse: 23.5593 - val_loss: 18.2555 - val_mse: 18.2555
Epoch 19/50
12/12 [=====] - 0s 7ms/step - loss: 19.2828 - mse: 19.2828 - val_loss: 17.9827 - val_mse: 17.9827
Epoch 20/50
12/12 [=====] - 0s 5ms/step - loss: 22.7105 - mse: 22.7105 - val_loss: 20.5292 - val_mse: 20.5292
Epoch 21/50
12/12 [=====] - 0s 7ms/step - loss: 22.5749 - mse: 22.5749 - val_loss: 18.6884 - val_mse: 18.6884
Epoch 22/50
12/12 [=====] - 0s 5ms/step - loss: 19.0899 - mse: 19.0899 - val_loss: 58.9656 - val_mse: 58.9656
Epoch 23/50
12/12 [=====] - 0s 6ms/step - loss: 22.1505 - mse: 22.1505 - val_loss: 17.7170 - val_mse: 17.7170
Epoch 24/50
12/12 [=====] - 0s 5ms/step - loss: 21.5239 - mse: 21.5239 - val_loss: 17.5552 - val_mse: 17.5552
Epoch 25/50
12/12 [=====] - 0s 5ms/step - loss: 18.7696 - mse: 18.7696 - val_loss: 33.4806 - val_mse: 33.4806
Epoch 26/50
12/12 [=====] - 0s 5ms/step - loss: 21.0288 - mse: 21.0288 - val_loss: 37.7337 - val_mse: 37.7337
Epoch 27/50
12/12 [=====] - 0s 5ms/step - loss: 19.4240 - mse: 19.4240 - val_loss: 27.9656 - val_mse: 27.9656
Epoch 28/50
12/12 [=====] - 0s 7ms/step - loss: 18.7796 - mse: 18.7796 - val_loss: 16.4296 - val_mse: 16.4296
Epoch 29/50
12/12 [=====] - 0s 6ms/step - loss: 19.5165 - mse: 19.5165 - val_loss: 36.3188 - val_mse: 36.3188

```

```
# Preview the mean value of training and validation data
y_train.mean(), y_val.mean()
```

```
(22.235537, 24.89756)
```

```
# Evaluate the model on the test data
print("Evaluation on Test data \n")
loss, mse = model.evaluate(X_test, y_test, batch_size=32)
print(f"\nModel loss on test set: {loss}")
print(f"Model mean squared error on test set: {(mse):.2f}")
```


```
Evaluation on Test data
```

```
4/4 [=====] - 0s 4ms/step - loss: 18.0088 - mse: 18.0088
```

```
Model loss on test set: 18.00882339477539
```

```
Model mean squared error on test set: 18.01
```

```
# Plot the loss curves
pd.DataFrame(history.history).plot(figsize=(6, 4), xlabel="Epochs", ylabel="Loss", title='Loss Curves')
plt.show()
```



```
# Make predictions
y_pred = model.predict(X_test)
```

```
# View the first prediction
y_pred[0]
```

```
4/4 [=====] - 0s 4ms/step
array([19.403996], dtype=float32)
```

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.